

SIMRAN KAUR

✉ skaur@princeton.edu 🌐 www.kaursim.com

EDUCATION

Princeton University

Ph.D. Candidate in Computer Science Department (Advisor: Sanjeev Arora).

Princeton, NJ
2022 - Present

Carnegie Mellon University

B.S. Artificial Intelligence, School of Computer Science.

Pittsburgh, PA
2018-2022

GPA: 3.95/4.00

HONORS

CMU Senior Leadership Recognition Recipient

May 2022

Phi Beta Kappa

May 2022

CMU SCS College Honors (successful completion of senior thesis)

May 2022

Dean's List

Spring 2019 – May 2022

PUBLICATIONS

[1] How does RL Post-Training Induce Skill Composition? A Case Study on Countdown [\[Link\]](#)

Simon Park*, **Simran Kaur***, Sanjeev Arora.

Spotlight at *Efficient Reasoning Workshop* (NeurIPS), 2025.

In *MATH-AI Workshop* (NeurIPS), 2025.

[2] Instruct-SkillMix: A Powerful Pipeline for LLM Instruction Tuning [\[Link\]](#)

Simran Kaur*, Simon Park*, Anirudh Goyal, Sanjeev Arora.

In *Proceedings of the 13th International Conference on Learning Representations (ICLR)*, 2025.

Oral at *Compositional Learning Workshop* (NeurIPS), 2024.

In at *Compositional Learning Workshop* (NeurIPS), 2024.

[3] Can Models Learn Skill Composition from Examples? [\[Link\]](#)

Haoyu Zhao, **Simran Kaur**, Dingli Yu, Anirudh Goyal, Sanjeev Arora.

In *Advances in Neural Information Processing Systems 37* (NeurIPS), 2024.

[4] Skill-Mix: a Flexible and Expandable Family of Evaluations for AI models [\[Link\]](#) [\[Quanta Article\]](#)

Dingli Yu, **Simran Kaur**, Arushi Gupta, Jonah Brown-Cohen, Anirudh Goyal, Sanjeev Arora.

In *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024.

In *Workshop on Distribution Shifts* (NeurIPS), 2023.

[5] Disentangling the Mechanisms behind Implicit Regularization in SGD [\[Link\]](#)

Zachary Novack, **Simran Kaur**, Tanya Marwah, Saurabh Garg, Zachary C. Lipton.

In *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023.

Spotlight in *Higher Order Optimization in Machine Learning Workshop* (NeurIPS), 2022. [Best Poster Award]

[6] On the Maximum Hessian Eigenvalue and Generalization [\[Link\]](#)

Simran Kaur, Jeremy Cohen, Zachary C. Lipton.

Contributed talk at *"I Can't Believe It's Not Better!" Workshop* (NeurIPS), 2022.

[7] Gradient Descent on Neural Networks Typically Occurs at the Edge of Stability [\[Link\]](#)

Jeremy M. Cohen, **Simran Kaur**, Yuanzhi Li, Zico Kolter, Ameet Talwalkar.

In *Proceedings of the 9th International Conference on Learning Representations (ICLR)*, 2021.

In *Opt2020: 12th Annual Workshop on Optimization for Machine Learning* (NeurIPS), 2020.

[8] Are Perceptually-Aligned Gradients a General Property of Robust Classifiers? [\[Link\]](#)

Simran Kaur, Jeremy Cohen, Zachary C. Lipton.

In *Science Meets Engineering of Deep Learning Workshop* (NeurIPS), 2019.

TEACHING

Princeton University

- Teaching Assistant for COS 324: *Introduction to Machine Learning*.
Fall 2023 (for Junior Research Work), Spring 2024.

Carnegie Mellon University

- Teaching Assistant for 15281 *Artificial Intelligence: Representation and Problem Solving*.
Spring 2020, Fall 2020, Spring 2021 (Head TA), Fall 2021 (Head TA), Spring 2022 (Head TA)
- Teaching Assistant for 10301/10601 *Introduction to Machine Learning* (Undergraduate and Graduate Level).
Summer 2020.

SERVICE

- Conference Reviewing: NeurIPS, ICLR 2023 – Present
- Co-organizer for Princeton Algorithms & Machine Learning (Alg-ML) Seminar [\[Link\]](#) Fall 2023 – Spring 2024
- Princeton Language and Intelligence (PLI) Blog Board Member [\[Link\]](#) Fall 2023 – Present

SKILLS

Programming, Frameworks & Softwares: Python, C, Java, Standard ML, R, LaTeX, PyTorch, Matlab, Jupyter Notebook, Git